

1) Which tasks have been completed?

I have done the initial data cleaning and tokenization for the collection of reviews. I also have an initial Naïve Bayes classifier trained to distinguish positive from negative reviews. I will further improve this to classify ratings rather than a binary positive/negative.

2) Which tasks are pending?

After finishing the classifier, the next step will be to set up a server which can take in a review and return a rating based on what the classifier predicts it to be. After that, I will work on the UI to interact with the server.

3) Are you facing any challenges?

Learning NLTK (Natural Language Toolkit) slowed me down quite a bit. I've also realized I need to cut down on data because run times are very slow if I try to use everything.